

Magnetism

The loop that doesn't close: why one consistency requirement forces every gauge force at once, and what a magnetic field actually is once you take that seriously.

v1.1 · 2026-08-10 (v1.0 2026-08-03; r2017 sync pass + octahedral-ablation note, see changelog) · intuition-first synthesis paper; register tags as in the sibling papers: MEASURED DERIVED CONDITIONAL OPEN

One-line claim. Magnetism is not a second force that happens to rhyme with electricity — it is what one and the same consistency requirement between observers looks like when you walk it around a closed loop and it doesn't come back to where it started. The requirement forces a single algebra containing every gauge force at once; electromagnetism is one thread of that algebra; and the magnetic field, specifically, is the holonomy — the leftover twist — of that thread's connection around a cycle of overlapping observers. Everything past that is bookkeeping.

1 · Two claims that keep getting run together

There is an old intuition, stated most famously by Walter Russell in 1926, that electricity and magnetism are not two forces but one force read in two directions — "the south wind and the north wind are not two winds." This project already audited that specific claim (The Universal One, Redrawn §5.2) and the verdict stands: as literal force-ontology it is wrong — electromagnetic radiation adds to gravity's pull rather than opposing it, and the true expansive pole is dark energy, not magnetism. As a structural rhyme — two opposed tendencies meeting at a balance point — it survives, vindicated in shape, in stellar hydrostatic equilibrium. This paper does not reopen that verdict.

What it does is something Russell had no access to: an actual derivation, inside Observer Patch Holography (OPH), of a real "one force" moment — not a poetic unity of electricity and magnetism specifically, but a forced unity of *every* gauge force's underlying symmetry, electromagnetism's included, from one requirement about how

observers agree with each other. Magnetism's place inside that requirement turns out to be sharp, specific, and different from anything Russell's diagram anticipated: not a force at all in the first instance, but a **consistency defect** — the amount by which walking a record around a loop fails to bring it home.

2 • One requirement, one algebra

OPH models an observer as a finite *patch*: a local state, a boundary with twelve ports arranged with the symmetry of an icosahedron, a readback record, and a set of repair moves for reconciling disagreements with neighboring patches. Two requirements are placed on this carrier, and nothing else:

- **A1 — complete, faithful, reversible response.** Each patch has a response space that reacts to everything at its boundary, and the reaction is invertible: nothing is thrown away.
- **A2 — endogenous agreement.** Whenever two overlapping patches reconcile their shared records, the reconciliation has to be implemented *from inside* the same response space each patch already has — not by appeal to an outside referee, an ambient clock, or an assumed symmetry group.

Given the icosahedral incidence of the twelve ports (the rotation group is A_5 , order sixty) and the requirement that only one direction in the response space stay fixed under it, these two requirements leave exactly one algebra standing:

$$\mathfrak{g} \cong \mathfrak{u}(1) \oplus \mathfrak{su}(2) \oplus \mathfrak{su}(3)$$

— the Standard Model's entire gauge algebra. No gauge group is fed in as a premise anywhere in A1 or A2; the classification is a finite, machine-checked theorem (Lean-verified) about what algebra a twelve-port response space with this symmetry and this transport rule can support at all. The centreless compact alternative of the same dimension, $\mathfrak{su}(2)^4$, is excluded outright — it cannot support the one-dimensional fixed direction the icosahedral action demands. DERIVED (Mueller, Osika, Ponder, Xue, Nguyen, Visser & Matscheko, *From Observer Consensus to Standard Physics*, release r2017, Theorem 8.2 "A1–A2 gauge Lie-type forcing" — the theorem also fixes the carrier decomposition: up to the outer automorphism of A_5 , the $\mathfrak{su}(2)$ ideal carries the

3 and the $\text{su}(3)$ ideal carries $3' \oplus 5$; companion detail in *Deriving Standard Model Gauge Structure from Observer Overlap Consistency*, release r2017)

Ablation (added v1.1, 2026-08-10). The icosahedral incidence in the hypothesis is load-bearing, not decorative, and this is now receipted by ablation. The plane admits exactly two self-similar lattice grammars; closing the other one (the square grammar) yields a screen with eight cube-corner ports and rotation group S_4 , and running the same A1+A2 machinery on it **fails closure entirely**: endogenous implementation fixes the centre of the response algebra pointwise and transitivity makes the fixed space one-dimensional, so the centre is at most one-dimensional — which leaves either a seven-dimensional compact semisimple algebra (none exists) or $\text{su}(3)$ alone (excluded by faithful port-equivariance). No admissible response algebra exists on the cube screen. Twelve icosahedral ports are not one option among many; they are the unique self-similar closure on which the forcing closes at all. DERIVED (receipts:

`lpoh/sims/PREREG_GRAMMAR_DISCRIMINATOR_2026-08-10.md` , GD-H and C4-discharge sections, with `runtime_receipts/gd_h_forcing_attempt.json` and `gd_h_c4_discharge.json` ; narrative home ashayana_deane.html §8.1b)

This is the real "one force" moment, and it is stronger than Russell's, not weaker. He wanted electricity and magnetism to be one thing. What OPH actually forces is that the strong force, the weak force, and electromagnetism's symmetry types are not three separate discoveries at all — they are one forced consequence of one requirement: that twelve-port observers can reconcile their records using only their own bookkeeping, with no outside referee.

3 • Picking out electricity and magnetism's own thread

The forced algebra has three pieces. Which one is electromagnetism's, specifically — as opposed to, say, the raw hypercharge direction it sits next to? That needs one more ingredient beyond A1 and A2: a symmetry-breaking attachment (a single scalar direction, playing the role the Higgs field plays in the Standard Model). Once that attachment is declared, the surviving unbroken generator is forced to be

$$Q = T_3 + Y,$$

and the corresponding factor $U(1)_Q$ is the electromagnetic one — distinct from the weak-charged carriers $W_{\pm}$ and the neutral weak carrier Z , which the same attachment produces alongside it. CONDITIONAL (on the declared one-Higgs branch — this step is not available from A1 and A2 alone)

4 • The field is just curvature — automatic, not extra physics

Once $U(1)_Q$ is realized as an actual connection A_Q on the observer network — a rule for comparing phases between overlapping patches along that one direction — its field strength is simply

$$F_Q = dA_Q.$$

This is not a new law of nature. It is what the word "curvature" means once you have a connection: the field is the connection's own rate of change, nothing added. DERIVED (falls out of having a connection at all; Corollary 6.6d, first clause)

5 • Where magnetism actually lives: the loop that doesn't close

Here is the part that matters most, and it is proved independently of the electromagnetic identification above — it is a general fact about *any* overlap-transport structure on the twelve-port network. Take a closed loop of overlapping patches — patch 1 overlaps patch 2, which overlaps patch 3, ... which overlaps patch 1 again. Transport a record around that loop, patch by patch. Nothing in A1 or A2 guarantees you get back the record you started with. The leftover discrepancy — the amount by which the loop fails to close — is a genuine, gauge-invariant quantity: a class in the network's Čech cohomology, and specifically, for two paths sharing the same endpoints, **the holonomy of the loop they form together.** DERIVED (the central-defect and residual-holonomy theorem, Theorem 6.12 in the same paper)

That is, letter for letter, what a magnetic field *is*, once you have a connection to hang it on. The Aharonov–Bohm phase a charged particle picks up going around a solenoid, the flux through a loop that Stokes' theorem converts into a line integral of the vector potential — these are not new facts about nature layered on top of the algebra; they are the same holonomy statement, in the language physics already uses for it. In OPH's

own vocabulary: a magnetic field is the shape of a consistency defect — the failure of two ways of transporting a record around a loop to agree with each other. Not a push on a particle in the first instance. A bookkeeping fact about closure.

This is not an analogy imported from outside. The same source's own treatment of the network's dynamics — a lattice gauge Hamiltonian on the "screen worldvolume," in its own words — names the smallest closed loop (one plaquette) the "*magnetic*" term, sitting beside the electric terms on links and the Gauss-law constraints at vertices. OPH already calls this object magnetism, in its own internal language, independently of anything this paper adds.

DERIVED

6 • Three layers, not two: ledger, transport, wave

An earlier pass at this question (this project's own attempt to derive a force law directly from OPH's bare edge-local records — the raw ledger of separation r and its own time-derivatives at one instant, carried at every port) found that layer could only ever write **Weber-class** laws: central, relational, depending only on $r, \dot{r}, \ddot{r}$ at each edge, with no way to compare a vector's direction between two separated points. That is provably too poor a structure to carry a genuine magnetic field — the derivation is explicit that the full Grassmann/Lorentz structure needs "transported vector comparisons" the bare ledger does not have. DERIVED

The resolution is not that OPH lacks the needed structure — it is that magnetism does not live in that layer at all. It lives one layer up, in the **overlap-transport groupoid** that A2 already requires: the rule by which two patches reconcile records is exactly a connection, and a connection is exactly rich enough to carry holonomy. Three layers, cleanly separated:

layer	what it can write	magnetism?
bare edge ledger ($r, \dot{r}, \ddot{r}$)	Weber-class central forces only	no — provably too poor
overlap-transport / connection (A2)	holonomy around closed loops	yes — this is its home

full propagating dynamics (Maxwell action)	retarded fields, radiation, force laws	the familiar wave-form of it, one premise further
---	---	--

This also reframes, rather than answers, the earlier bench question this project logged as an open obligation (OB-2, the record-layer-vs-wave-channel discriminator on an open hairpin circuit, 6.91 gf if ledger-direct against 5.35 gf if wave-mediated). That framing posed the choice as binary — ledger or wave. The transport-holonomy layer is a genuine third option sitting between them, and it is the one this derivation actually uses. Whether the hairpin discriminator still cleanly separates "ledger-direct" from "wave-mediated" once a transport layer is on the table is not resolved here and is worth a second look. [OPEN](#)

Addendum, 2026-08-03 — the second look, taken. Cross-checked against this project's own separate, already-run attempt to add a dynamics layer to the bare ledger (OPH_PHYSICAL_ATTACHMENT_2026-08-01.md §4, gate E1). Three findings, in brief — full derivation in OB-2's amendment: (i) two unrelated methods (E1's record-lag dynamics; this paper's algebraic gauge-forcing) independently confirm the bare ledger cannot carry magnetism alone — real, if modest, corroboration; (ii) a tempting tiebreaker — "the record branch is non-causal, the connection branch respects the light cone, so OPH's own Lorentz-kinematics branch should exclude it" — does not survive contact with this project's own no-signaling theorem (OB-1 T1): E1's record-lag mechanism already respects the same light-cone bound, so causality cannot be cited against it; (iii) both tracks reduce to one unified open question in two vocabularies — does the forced connection necessarily carry *independent radiative dynamics* (a real field, à la the Maxwell- action premise of §7) or can it be non-radiative (a two-body potential with no independent field entity, which is what classical Weber theory and the local attenuation branch both are)? Neither track supplies the missing premise. **OB-2 stays open** — sharper, not resolved. [OPEN](#)

Correction, same day. The addendum above was written from the papers' prose alone. Checked against the actual instrument (muellerberndt/oph-physics-sim) and it already runs this exact test: oph_fpe/gauge/kinetic_selector_sweep.py asks, verbatim, whether positivity, locality, and A5 covariance select a unique kinetic response law. Run directly: they do not — twelve candidate laws swept, a constructive two-law counterexample, receipts UNIQUE_GAUGE_KINETIC_SOURCE_RAY_RECEIPT: false and PHYSICAL_GAUGE_KINETIC_ACTION_RECEIPT: false . So no new theorem needs inventing — the relevant one has already been run, by Mueller's own instrument, and returned open for a

named reason (no derived physical gauge kinetic action). The sweep's own claim boundary names the next concrete step: a richer source dynamics might select a law. The same repo has a candidate for that — `oph_fpe/dynamics/canonical_seam_repair.py`, the real repair engine — not yet pointed at this question. Full detail in OB-2's second amendment. [OPEN](#)

Ran it. `canonical_seam_repair.py`'s own certificate confirms the repair generator is uniquely $-L$ up to a time-scale, with a forced ratio $(5-\sqrt{5} : 6 : 5+\sqrt{5})$ across the three nonconstant bands. Of the twelve candidate Hessians `kinetic_selector_sweep.py` already validated, exactly two (`edge_repair`, `double_edge_repair`) reproduce that same ratio beyond their onsite term — the other ten do not. A real narrowing ($12 \rightarrow$ a one-parameter scale family), composed entirely from two already-run certificates. **But the bridge premise — that the response Hessian should match the repair generator — is mine, not an OPH theorem, and even granted it says nothing about radiative-vs-non-radiative,** which stays Corollary 6.6d's untouched gap. Flagged, not pursued further for now (operator instruction). [OPEN](#)

7 • The honest boundary: what's forced, and what's assumed

Sections 2, 4, and 5 are the derived kinematic core: one algebra, forced; a connection's curvature, automatic; holonomy, the general shape of a magnetic observable. None of that is yet Maxwell's equations. To get there, one further premise has to be declared — the Maxwell action itself,

$$S = -1/(2g^2) \int F_Q \wedge *F_Q + \int A_Q \wedge *J_Q,$$

after which the Euler–Lagrange equations give $dF_Q = 0$ (the no-monopole identity, automatic anyway once $F_Q = dA_Q$) and $d*F_Q = g_Q^2 *J_Q$ (Gauss's law and the Ampère–Maxwell law, sourced by the current). After the usual normalization, these are Maxwell's equations. The source paper's own words are the right ones to keep here: *"The compact-group reconstruction and connection label do not themselves supply this action or its nonzero kinetic coefficient."* [CONDITIONAL](#) (Corollary 6.6d)

So the honest summary has two halves, and they should not be blurred into one. OPH forces the *existence* of the right kind of symmetry for electromagnetism, forces the connection once you point at that symmetry's abelian piece, and forces holonomy as the

general shape of what a magnetic field is. It does not, on its own, force the specific dynamical law that makes real electromagnetism propagate the way it does, or fix the numerical size of the coupling (that number belongs to a separate closure — the paper's screen-grain proposal reads back the fine-structure constant to within a few parts in 10^{-6} of the 2022 CODATA value, but is tagged there as a diagnostic against a declared source map, not a derivation of α from A1–A2). CONDITIONAL




Figure 1 — the loop that doesn't close. A record is transported around five overlapping patches. Nothing in A1 or A2 forces the trip to end where it began; the discrepancy between "go the long way" and "go the short way" is the loop's holonomy — a gauge-invariant number, and (once the electromagnetic branch is realized) exactly the magnetic flux through that loop.

8 • Back to Russell, once, correctly

It is worth being precise about what this derivation does and does not hand back to the intuition that started this inquiry. Russell wanted electricity and magnetism to be the same push in opposite directions — the audited verdict on that specific claim stands unchanged (The Universal One, Redrawn §5.2): wrong as force-ontology, vindicated only as a structural shape. What OPH actually supplies is a different, sharper unity: not electricity and magnetism as mirror pushes, but *one consistency requirement between observers, forcing one algebra that contains every gauge charge at once* — and magnetism, inside that algebra, is what shows up when the requirement is read around a cycle instead of along a line. That is a better sentence than his, and it costs nothing to say plainly that it took a century, twelve ports, and two axioms to get there.

Claim boundary. This paper introduces no new OPH theorem — every derived step is sourced to the external Observer Patch Holography papers cited below, and every conditional step is marked exactly where its source marks it. What is new here is the synthesis and the intuitive reading: magnetism-as-holonomy (§5), the three-layer ledger/transport/wave taxonomy (§6), and the reframing of the Weber-vs-wave bench question. Do not cite this paper as a source for the gauge-forcing theorem, the curvature identity, or the holonomy theorem — cite the originals. Do not read §6's reframing as a resolution of OB-2; it is a sharpened question, not a closed one. Do not treat the fine-structure reference in §7 as an OPH derivation of α — the source paper itself tags it a diagnostic. This paper's own local `oph_programme.html / shape_of_the_universe.html` dialect predates and has not yet absorbed the axioms used here; a separate sync pass is owed and is not this paper's job.

Sources. This paper is built directly on the external Observer Patch Holography programme (Pragma Research Inc., B. Mueller corresponding author; A. Osika among the co-authors), not the local project's earlier `oph_programme.html` dialect:

- Mueller, Osika, Poneder, Xue, Cassie, Nguyen, Visser, Anirudha, Matscheko, Hill & Glynn (2026). *From Observer Consensus to Standard Physics: Normal Forms, the Standard Model Lie Type, and a Route to the Einstein Field Equation*. Release r2017 (2026-08-06; this paper first cited the r2013 release — sync pass 2026-08-10, all cited claims verified against r2017). The source of Theorem 8.2 (p. 22, "A1–A2 gauge Lie-type forcing").
- Mueller, Osika, Poneder, Xue, Nguyen, Visser & Matscheko (2026). *Deriving Standard Model Gauge Structure from Observer Overlap Consistency*. Release r2017. The source of Corollaries 6.6c–d (pp. 65–66) and Theorem 6.12 (pp. 70–71) as cited above — numbering unchanged from the release first cited.
- Mueller, Osika, Poneder, Xue, Nguyen, Visser & Matscheko (2026). *Recovering Observer Spacetime and Einstein Dynamics from Overlap Consistency*.

All public at `github.com/FloatingPragma/observer-patch-holography` (which also hosts the **public Lean development** — the finite Lie-type classification arithmetic behind Theorem 8.2 is formalized sorry-free in `Lean/Screen/A50PH.lean`, including the $11 = 3+8$ and $12 = 3+3+3+3$ uniqueness lemmas) and rendered at `muellerberndt.github.io/oph/papers/`.

Changelog. v1.0 · 2026-08-03 · initial version, citing the r2013 releases. v1.1 · 2026-08-10 · **r2017 sync pass:** every cited claim re-verified against the r2017 releases — Theorem 8.2 re-anchored to its actual home (*From Observer Consensus to Standard Physics*, p. 22; its statement now also fixes the carrier decomposition quoted in §2); Corollaries 6.6c–d and Theorem 6.12 confirmed unchanged in the gauge paper; sources box updated with page anchors, the rendered-site URL, and the public Lean development

(`Lean/Screen/A50PH.lean` , sorry-free). Added the §2 octahedral-ablation note: the same A1+A2 machinery on the square-grammar (cube) closure fails entirely — icosahedral incidence is the unique self-similar closure on which the forcing closes (receipts in `lpoh/sims/` and `runtime_receipts/`).

Alexander Osika · August 2026 · sibling papers: Light · The Electron · The Shape of the Universe · [The Shape of Gravity](#) · The Universal One, Redrawn